

bayespca Package

Davide Vidotto d.vidotto@uvt.nl

2019-05-31

bayespca: A package for Variational Bayes PCA

Theoretical background

Principal Components Analysis (PCA) allows performing dimensionality reduction via matrix factorization. While there are several ways to express a PCA model, in what follows will we consider the formulation

$$X = XWP^T + E,$$

where X is a $I \times J$ data matrix (I is the number of units; J the number of continuous variables); W is a $J \times D$ weight matrix ($D \leq J$ is the rank of the reduced matrix); P is the orthogonal loading matrix, such that $P^T P = I_{D \times D}$; and E is an $I \times J$ error matrix. The D principal components can be retrieved with $Z = XW$. In this context, the focus of the inference is typically on W . In particular, when J is large and the main inferential goal is components' interpretation, it is important for the analyst to obtain simple and interpretable components.

The **bayespca** package allows performing the following operations:

1. estimation of the PCA model, with a Variational Bayes algorithm;
2. regularization of the elements of W by means of its prior variances;
3. variable selection, via a Stochastic Search Variable Selection method (a form of “spike-and-slab” prior).

The Variational Bayes algorithm sees the columns of W as latent variables, and P as a fixed parameter. Furthermore, the residuals E are assumed to be distributed according to a Normal distribution with mean 0 and variance σ^2 . The following prior is assumed for the d -th column of W :

$$w_d \sim MVN(0, T_d)$$

where $MVN()$ denotes the density of the Multivariate Normal Matrix, and T_d denotes the prior (diagonal) covariance matrix of the d -th component. The j -th element of the diagonal of T_d will be denoted τ_{dj} . where $MVN()$ denotes the density of the Multivariate Normal Matrix, and T_d denotes the prior (diagonal) covariance matrix of the d -th component. The j -th element of the diagonal of T_d will be denoted τ_j .

The bayespca package

Variational Bayes PCA is implemented through the **vbpca** function, which takes the following arguments as inputs:

- **X** the input matrix;
- **D** the number of components to be estimated;
- **maxIter** the maximum number of iterations for the Variational Bayes algorithm;
- **tolerance** convergence criterion of the algorithm (relative difference between ELBO values);
- **verbose** logical parameter which prints estimation information on screen when TRUE;
- **tau** value of the prior variances; starting value when **updatetau=TRUE** or **priorvar!='fixed'**
- **updatetau** logical parameter denoting whether the prior variances should be updated when **priorvar='fixed'**;

- `priorvar` character argument denoting whether the prior variances should be 'fixed', or random with 'jeffrey' or 'invgamma' priors;
- `SVS` logical argument which activates Stochastic Variable Selection when set to TRUE;
- `priorInclusion` prior inclusion probabilities for the elements of W in the model;
- `global.var` logical parameter which activates component-specific prior variances when set to TRUE;
- `control` other control parameters, such as Inverse Gamma hyperparameters (see `?vbpc_control` for more information).

`vbpc` returns a `vbpc` object, which is a list containing various aspect of the model results. See `?vbpc` for further information. Internally, `vbpc` calls a C++ function (written with Rcpp) to estimate the model.

In what follows, the various estimation modalities allowed by `vbpc` will be introduced. For presentation purposes, a synthetic data matrix with $I = 100$ rows and $J = 20$ columns generated from three components will be used:

```
set.seed(141)
I <- 100
J <- 20
V1 <- rnorm(I, 0, 50)
V2 <- rnorm(I, 0, 30)
V3 <- rnorm(I, 0, 10)
X <- matrix(c(rep(V1, 7), rep(V2, 7), rep(V3, 6)), I, J)
X <- X + matrix(rnorm(I * J, 0, 1), I, J)
```

I will now proceed with the estimation of the PCA model.

Levels of regularization on the W matrix

Fixed tau

With fixed tau, it is possible to specify the model as follows:

```
# Install and load package
# devtools::install_github("davidedvt/bayespca")
library(bayespca)

# De-activate data center and scaling;
ctrl <- vbpc_control(center = FALSE, scalecorrection = -1,
                     plot.lowerbound = FALSE)

# Estimate vbpc with fixed prior variances (equal to 1)
# for the elements of W
mod1 <- vbpc(X, D = 3, maxIter = 1e+03, priorvar = 'fixed',
             control = ctrl, verbose = FALSE )

# Test the class of mod1:
is.vbpc(mod1)
```

```
## [1] TRUE
```

The estimate posterior means of the W matrix can be viewed with:

```
mod1$muW
```

```
##           Component 1 Component 2   Component 3
## [1,] -0.376589700 -0.04416511  0.0003399127
## [2,] -0.373939778 -0.04582346 -0.0111489595
## [3,] -0.375148658 -0.04305857 -0.0078831845
## [4,] -0.374770078 -0.04473100 -0.0031124941
## [5,] -0.376808027 -0.04285792 -0.0100250681
## [6,] -0.375114066 -0.04446329 -0.0015012124
## [7,] -0.375069347 -0.04364081 -0.0007181138
## [8,]  0.043916074 -0.37610685 -0.0194987844
## [9,]  0.044338996 -0.37382690 -0.0224165536
## [10,] 0.043216238 -0.37319456 -0.0161965576
## [11,] 0.043432789 -0.37311089 -0.0246530518
## [12,] 0.045420158 -0.37574267 -0.0200072059
## [13,] 0.045158091 -0.37616395 -0.0206149568
## [14,] 0.044605651 -0.37571347 -0.0144837533
## [15,] 0.002905219  0.02229238 -0.4057459837
## [16,] 0.003409761  0.02199152 -0.4068881894
## [17,] 0.003232845  0.02063894 -0.4106993908
## [18,] 0.002919709  0.02319335 -0.4056785190
## [19,] 0.002019259  0.02192116 -0.4088024260
## [20,] 0.001874207  0.02043128 -0.4078308025
```

and the P matrix:

```
mod1$P
```

```
##           Component 1 Component 2   Component 3
## [1,] -0.376589904 -0.04416517  0.0003399179
## [2,] -0.373939981 -0.04582353 -0.0111491289
## [3,] -0.375148862 -0.04305863 -0.0078833043
## [4,] -0.374770282 -0.04473106 -0.0031125414
## [5,] -0.376808232 -0.04285797 -0.0100252205
## [6,] -0.375114270 -0.04446335 -0.0015012352
## [7,] -0.375069551 -0.04364087 -0.0007181247
## [8,]  0.043916097 -0.37610735 -0.0194990808
## [9,]  0.044339020 -0.37382740 -0.0224168943
## [10,] 0.043216262 -0.37319506 -0.0161968038
## [11,] 0.043432813 -0.37311139 -0.0246534264
## [12,] 0.045420183 -0.37574317 -0.0200075099
## [13,] 0.045158115 -0.37616446 -0.0206152700
## [14,] 0.044605675 -0.37571398 -0.0144839734
## [15,] 0.002905220  0.02229241 -0.4057521497
## [16,] 0.003409762  0.02199155 -0.4068943728
## [17,] 0.003232846  0.02063897 -0.4107056321
## [18,] 0.002919710  0.02319338 -0.4056846840
## [19,] 0.002019260  0.02192119 -0.4088086384
## [20,] 0.001874208  0.02043131 -0.4078370001
```

Among other things, the function returns also the model evidence lower bound (ELBO) and the estimation time:

```
mod1$elbo
```

```
## [1] -2834.277
```

```
mod1$time
```

```
##      user  system elapsed
##         0         0         0
```

Fixed, updatable tau

The prior variances τ_{dj} can also be updated via Type-II Maximum Likelihood (empirical Bayes updates):

```
mod2 <- vbPCA(X, D = 3, maxIter = 1e+03, priorvar = 'fixed',
              updatetau = TRUE, control = ctrl, verbose = FALSE )
```

```
mod2$muW
```

```
##      Component 1 Component 2 Component 3
## [1,] -3.774720e-01 -0.051470102 -0.001822156
## [2,] -3.744848e-01 -0.025307896 -0.002337156
## [3,] -3.747070e-01 -0.039179989 -0.002345241
## [4,] -3.697882e-01 -0.062258505 -0.002242303
## [5,] -3.794353e-01 -0.031311390 -0.002225193
## [6,] -3.811333e-01 -0.058594825 -0.001901847
## [7,] -3.709497e-01 -0.037084346 -0.001895318
## [8,]  4.572526e-02 -0.388171353 -0.019796029
## [9,]  4.119102e-02 -0.376636645 -0.023267758
## [10,] 2.518537e-02 -0.376940208 -0.006967316
## [11,] 5.297341e-02 -0.374414696 -0.022267903
## [12,] 4.534880e-02 -0.374381287 -0.013986121
## [13,] 6.385742e-02 -0.370055099 -0.031673499
## [14,] 3.111306e-02 -0.364729611 -0.002535536
## [15,] 1.288107e-05  0.006270320 -0.406011655
## [16,] 1.305259e-05  0.018689756 -0.407077598
## [17,] 1.211560e-05  0.009078483 -0.411159927
## [18,] 1.039548e-05  0.036022360 -0.398788351
## [19,] 1.331297e-05  0.005889322 -0.410538009
## [20,] 1.326763e-05  0.035028974 -0.412758321
```

The matrix of the inverse prior variances can be called with

```
mod2$invTau
```

```
##      Component 1 Component 2 Component 3
## [1,] 6.711653e+00 185.292814 30792.082987
## [2,] 6.811715e+00 435.456531 24282.641778
## [3,] 6.733579e+00 246.605678 24174.268756
## [4,] 6.906649e+00 137.269136 25240.818705
## [5,] 6.630610e+00 334.561358 25606.215182
## [6,] 6.535962e+00 151.540810 29303.191067
## [7,] 6.916642e+00 278.901798 29505.931819
## [8,] 2.000395e+02  6.283433  670.346568
## [9,] 2.297542e+02  6.662857  568.834746
## [10,] 4.206301e+02  6.647409 1984.202810
## [11,] 1.678142e+02  6.740676  596.809053
## [12,] 2.033657e+02  6.710938  964.165004
## [13,] 1.294742e+02  6.828244  390.109550
## [14,] 3.541861e+02  7.207997 5635.426286
```

```
## [15,] 3.077602e+05 2373.138269      5.837005
## [16,] 3.081676e+05  748.154670      5.771306
## [17,] 3.180735e+05 1556.626018      5.676298
## [18,] 3.115912e+05  349.915331      5.997557
## [19,] 3.116412e+05 2503.322334      5.710167
## [20,] 3.100489e+05  362.448139      5.591906
```

Random tau: Jeffrey's prior

By assuming Jeffrey's hyperpriors on $\tau_{d,j}$ we set:

$$p(\tau_{d,j}) \propto \frac{1}{\tau_{d,j}}.$$

The following code runs the algorithm with Jeffrey's priors on tau:

```
mod3 <- vbpc(X, D = 3, maxIter = 1e+03,
             priorvar = 'jeffrey', control = ctrl, verbose = FALSE )
```

```
mod3$muW
```

```
##           Component 1 Component 2 Component 3
## [1,] -3.762072e-01 -0.051669847 -0.001718286
## [2,] -3.753663e-01 -0.024957643 -0.002179155
## [3,] -3.755221e-01 -0.039050104 -0.002184467
## [4,] -3.694122e-01 -0.062405832 -0.002095490
## [5,] -3.804778e-01 -0.031175496 -0.002080751
## [6,] -3.804682e-01 -0.058826990 -0.001785660
## [7,] -3.705392e-01 -0.036933654 -0.001783597
## [8,]  4.681938e-02 -0.388318730 -0.019714801
## [9,]  4.140105e-02 -0.376744554 -0.023221555
## [10,] 2.630865e-02 -0.376993005 -0.006755614
## [11,] 5.078766e-02 -0.374364439 -0.022225937
## [12,] 4.534032e-02 -0.374417607 -0.013628847
## [13,] 6.314263e-02 -0.370025997 -0.031527208
## [14,] 3.138958e-02 -0.364536417 -0.002429599
## [15,] -5.988475e-06  0.006004422 -0.406018863
## [16,] -5.753386e-06  0.018302788 -0.406955301
## [17,] -6.315337e-06  0.008713880 -0.411158474
## [18,] -7.980688e-06  0.036268370 -0.398755119
## [19,] -5.552139e-06  0.005629840 -0.410622420
## [20,] -5.662205e-06  0.035039733 -0.412865313
```

```
mod3$invTau
```

```
##           Component 1 Component 2 Component 3
## [1,] 6.755108e+00  184.397432 34474.447450
## [2,] 6.781195e+00  442.648244 27480.918786
## [3,] 6.705945e+00  247.645663 27392.152386
## [4,] 6.919912e+00  136.854417 28507.875401
## [5,] 6.595924e+00  336.309756 28900.632235
## [6,] 6.557748e+00  150.778136 32956.292021
## [7,] 6.931301e+00  280.357630 33103.785341
## [8,] 1.941867e+02   6.278975  675.700374
```

```
## [9,] 2.284549e+02    6.659287    571.930720
## [10,] 4.008565e+02    6.645666   2054.421916
## [11,] 1.770003e+02    6.742445    599.980458
## [12,] 2.033806e+02    6.709787    993.678825
## [13,] 1.313076e+02    6.829335    393.545372
## [14,] 3.504575e+02    7.215396   5895.969600
## [15,] 3.431666e+05   2488.239113     5.836825
## [16,] 3.436743e+05    768.309909     5.774690
## [17,] 3.546114e+05   1628.566791     5.676387
## [18,] 3.471180e+05    348.427445     5.998570
## [19,] 3.475540e+05   2628.776461     5.707931
## [20,] 3.459061e+05    363.790032     5.589196
```

Random tau: Inverse Gamma prior

It is possible to specify an inverse gamma prior on $\tau_{d,j}$:

$$\tau_{d,j} \sim IG(\alpha, \beta)$$

with α shape parameter and β scale parameter. The following code implements an $IG(2, .5)$ prior on the variances:

```
# Set hyperparameter values
ctrl2 <- vbPCA_control(center = FALSE, scalecorrection = -1,
                      plot.lowerbound = FALSE,
                      alphatau = 2, betatau = .5)

# Estimate the model
mod4 <- vbPCA(X, D = 3, maxIter = 1e+03, priorvar = 'invgamma',
             control = ctrl2, verbose = FALSE )

mod4$muW
```

```
##      Component 1 Component 2 Component 3
## [1,] -0.376590725 -0.04416826  0.0002947089
## [2,] -0.373916478 -0.04580818 -0.0111139002
## [3,] -0.375152260 -0.04306330 -0.0078485203
## [4,] -0.374771451 -0.04473662 -0.0031287802
## [5,] -0.376819149 -0.04285298 -0.0099904658
## [6,] -0.375128188 -0.04445803 -0.0015201459
## [7,] -0.375056611 -0.04365135 -0.0007401483
## [8,]  0.043918292 -0.37612390 -0.0195191684
## [9,]  0.044336996 -0.37380699 -0.0224249101
## [10,] 0.043215418 -0.37316141 -0.0162245050
## [11,] 0.043435548 -0.37309952 -0.0245960805
## [12,] 0.045416539 -0.37575273 -0.0200103648
## [13,] 0.045161055 -0.37621087 -0.0206044028
## [14,] 0.044603008 -0.37569186 -0.0144838955
## [15,] 0.002901159  0.02228837 -0.4056891063
## [16,] 0.003405260  0.02198287 -0.4068836504
## [17,] 0.003228832  0.02064976 -0.4107113982
```

```
## [18,] 0.002920806 0.02319222 -0.4056322666
## [19,] 0.002018556 0.02191743 -0.4087810436
## [20,] 0.001885825 0.02043591 -0.4078295957
```

```
mod4$invTau
```

```
##      Component 1 Component 2 Component 3
## [1,] 4.349798    4.952591    4.962213
## [2,] 4.357381    4.951847    4.961554
## [3,] 4.349198    4.947077    4.955868
## [4,] 4.347326    4.942587    4.952373
## [5,] 4.346511    4.949795    4.958338
## [6,] 4.348754    4.945842    4.955514
## [7,] 4.353855    4.952447    4.961810
## [8,] 4.940550    4.341632    4.948114
## [9,] 4.944064    4.351079    4.951230
## [10,] 4.946698    4.354576    4.954566
## [11,] 4.941841    4.351048    4.948103
## [12,] 4.940429    4.343068    4.948542
## [13,] 4.934790    4.337277    4.942629
## [14,] 4.954153    4.353743    4.962926
## [15,] 4.964065    4.961664    4.266733
## [16,] 4.955381    4.953028    4.256713
## [17,] 4.956688    4.954630    4.246359
## [18,] 4.954427    4.951810    4.259675
## [19,] 4.962441    4.960082    4.256343
## [20,] 4.950618    4.948561    4.250325
```

α and β can also be specified as D -dimensional array, in which case the Inverse Gamma will have component-specific hyperparameters:

$$\tau_{d,j} \sim IG(\alpha_d, \beta_d)$$

```
# Set hyperparameter values
```

```
ctrl3 <- vbPCA_control(center = FALSE, scalecorrection = -1,
                        plot.lowerbound = FALSE,
                        alphatau = c(.5, 50, 3), betatau = c(.5, .01, 10),
                        hypertype = 'component')
```

```
# Estimate the model
```

```
mod5 <- vbPCA(X, D = 3, maxIter = 1e+03, priorvar = 'invgamma',
              control = ctrl3, verbose = FALSE )
```

```
mod5$muW
```

```
##      Component 1 Component 2 Component 3
## [1,] -0.378535069 -0.022550177 0.0025320655
## [2,] -0.376010408 -0.021331002 -0.0088680201
## [3,] -0.377058121 -0.022705524 -0.0057334841
## [4,] -0.376765736 -0.022799099 -0.0008880582
## [5,] -0.378711737 -0.022502952 -0.0078832752
## [6,] -0.377093814 -0.019646614 0.0006966940
## [7,] -0.376981673 -0.026435232 0.0014645273
```

```
## [8,] 0.021618521 -0.045787586 -0.0012802351
## [9,] 0.022165897 -0.050285073 -0.0042910494
## [10,] 0.021098823 -0.037934476 0.0018542188
## [11,] 0.021297952 -0.049321093 -0.0065615922
## [12,] 0.023138118 -0.049443570 -0.0017949202
## [13,] 0.022842413 -2.348222240 -0.0025229191
## [14,] 0.022345189 -0.049142867 0.0037198670
## [15,] 0.002992116 0.002091087 -0.4063507109
## [16,] 0.003475947 -0.001948958 -0.4074613372
## [17,] 0.003205251 0.009070114 -0.4112473691
## [18,] 0.003062334 0.006593144 -0.4063446402
## [19,] 0.002077795 0.003818705 -0.4093929706
## [20,] 0.001857368 -0.003947850 -0.4083201743
```

```
mod5$invTau
```

```
##      Component 1 Component 2 Component 3
## [1,] 1.735068 4883.95439 0.3498320
## [2,] 1.737956 4896.56336 0.3498312
## [3,] 1.734465 4881.98803 0.3498046
## [4,] 1.733223 4881.05886 0.3497852
## [5,] 1.733528 4884.42780 0.3498147
## [6,] 1.734107 4913.05908 0.3498008
## [7,] 1.736710 4838.95652 0.3498308
## [8,] 1.974312 4527.17251 0.3497740
## [9,] 1.976205 4440.53154 0.3497924
## [10,] 1.977395 4665.76589 0.3498030
## [11,] 1.974916 4459.48648 0.3497789
## [12,] 1.974487 4457.06148 0.3497771
## [13,] 1.971541 18.24647 0.3497487
## [14,] 1.981271 4463.32575 0.3498394
## [15,] 1.982402 5007.31481 0.3469784
## [16,] 1.977961 5007.38609 0.3469217
## [17,] 1.978586 4988.06776 0.3468739
## [18,] 1.977499 4997.44534 0.3469332
## [19,] 1.981517 5004.88288 0.3469270
## [20,] 1.975535 5004.38616 0.3468870
```

Notice the different level of regularization obtained across the different components. In order to activate these ‘component-specific’ hyperpriors, `hypertype = 'component'` was specified.

Random tau, random betatau

It is also possible to specify a Gamma hyperprior on β (while α remains fixed):

$$\beta \sim Ga(\gamma, \delta).$$

This is achievable by setting `gammatau` (and `deltatau`) larger than 0 in the control parameters:

```
# Specify component-specific Gamma(.01, 10) hyperpriors on betatau
ctrl4 <- vbpc_control(center = FALSE, scalecorrection = -1,
                      plot.lowerbound = FALSE,
                      alphatau = 1, betatau = 1,
                      gammatau = .01, deltau = 10,
                      hypertype = 'component')
```



```
# Estimate the model
mod6 <- vbpc(X, D = 3, maxIter = 1e+03, priorvar = 'invgamma',
             control = ctrl14, verbose = FALSE )
```

```
mod6$muW
```

```
##      Component 1 Component 2 Component 3
## [1,] -0.376612836 -0.04414791 -0.001248665
## [2,] -0.373527098 -0.04527330 -0.009492228
## [3,] -0.375300115 -0.04330108 -0.006689935
## [4,] -0.374545218 -0.04475906 -0.003778637
## [5,] -0.377133411 -0.04290093 -0.008534217
## [6,] -0.375467993 -0.04449136 -0.002517447
## [7,] -0.374843813 -0.04378361 -0.001752866
## [8,]  0.044090921 -0.37645722 -0.020075378
## [9,]  0.044279637 -0.37340637 -0.022183245
## [10,] 0.043230956 -0.37276205 -0.017350634
## [11,] 0.043633248 -0.37285565 -0.022662370
## [12,] 0.045225801 -0.37597165 -0.020153747
## [13,] 0.045216689 -0.37689952 -0.020320873
## [14,] 0.044313381 -0.37547626 -0.014958361
## [15,] 0.002830475  0.02204270 -0.405124312
## [16,] 0.003209544  0.02187185 -0.407342357
## [17,] 0.003155435  0.02095304 -0.410906971
## [18,] 0.002824749  0.02288630 -0.405233175
## [19,] 0.002136077  0.02174539 -0.408749536
## [20,] 0.002176061  0.02077001 -0.407947441
```

```
mod6$invTau
```

```
##      Component 1 Component 2 Component 3
## [1,]  14.91448   50.12038   65.06736
## [2,]  15.08009   50.01542   64.75625
## [3,]  14.89640   49.61854   64.01402
## [4,]  14.91646   49.35990   64.02027
## [5,]  14.84777   50.00467   64.60364
## [6,]  14.89316   49.61383   64.18919
## [7,]  14.99621   50.11334   64.83842
## [8,]  49.11895   14.79168   63.03510
## [9,]  49.33578   14.99044   63.23453
## [10,] 49.58424   15.04221   63.72938
## [11,] 49.21586   15.00398   62.95653
## [12,] 49.08557   14.80959   63.03858
## [13,] 48.70446   14.70323   62.43163
## [14,] 50.23753   15.02297   64.98818
## [15,] 51.91718   51.58330   14.15424
## [16,] 51.29932   50.86616   13.95682
## [17,] 51.45181   51.11369   13.78557
## [18,] 51.19581   50.75885   14.05424
## [19,] 51.89061   51.52052   13.95095
## [20,] 50.91504   50.55465   13.87395
```

The posterior means of β can be accessed via

```
mod6$priorBeta
```

```
##           [,1]      [,2]      [,3]
## [1,] 0.02635365 0.02630866 0.02068517
## attr("names")
## [1] "beta 1" "beta 2" "beta 3"
```

hypertype specify the type of hyperprior for **beta**:

- 'common' implies $\beta \sim Ga(\alpha, \beta)$;
- 'component' implies $\beta_d \sim Ga(\alpha_d, \beta_d)$;
- 'local' implies $\beta_{dj} \sim Ga(\alpha_{dj}, \beta_{dj})$.

Similar to **alphatau** and **betatau**, **gammatau** and **deltatau** can also be D -dimensional arrays for component-specific hyperpriors on β .

Global prior variances

So far, the parameter **global.var** has always been set to **FALSE**, implying

$$w_{j,d} \sim N(0, \tau_{j,d}).$$

Setting **global.var = TRUE** will modify this formulation, which will switch to

$$w_{j,d} \sim N(0, \tau_d)$$

that is, component-specific variances (called ‘global variances’ in **vbpca**) will be estimated instead:

```
# Fixed prior global variances, updated via Type-II maximum likelihood:
mod7 <- vbpca(X, D = 3, maxIter = 1e+03, priorvar = 'fixed',
             updatetau = TRUE, control = ctrl, verbose = FALSE,
             global.var = TRUE)
```

```
mod7$muW
```

```
##      Component 1 Component 2 Component 3
## [1,] -0.376586376 -0.04416415  0.0003398288
## [2,] -0.373936478 -0.04582247 -0.0111462062
## [3,] -0.375145347 -0.04305764 -0.0078812377
## [4,] -0.374766771 -0.04473003 -0.0031117254
## [5,] -0.376804702 -0.04285698 -0.0100225924
## [6,] -0.375110756 -0.04446232 -0.0015008417
## [7,] -0.375066036 -0.04363986 -0.0007179364
## [8,]  0.043915686 -0.37609866 -0.0194939691
## [9,]  0.044338605 -0.37381876 -0.0224110177
## [10,] 0.043215857 -0.37318643 -0.0161925578
## [11,] 0.043432406 -0.37310277 -0.0246469635
## [12,] 0.045419757 -0.37573449 -0.0200022650
## [13,] 0.045157692 -0.37615576 -0.0206098658
## [14,] 0.044605257 -0.37570529 -0.0144801764
## [15,] 0.002905193  0.02229190 -0.4056457820
## [16,] 0.003409731  0.02199105 -0.4067877056
## [17,] 0.003232816  0.02063849 -0.4105979658
## [18,] 0.002919683  0.02319284 -0.4055783339
## [19,] 0.002019241  0.02192068 -0.4087014694
## [20,] 0.001874190  0.02043084 -0.4077300859
```

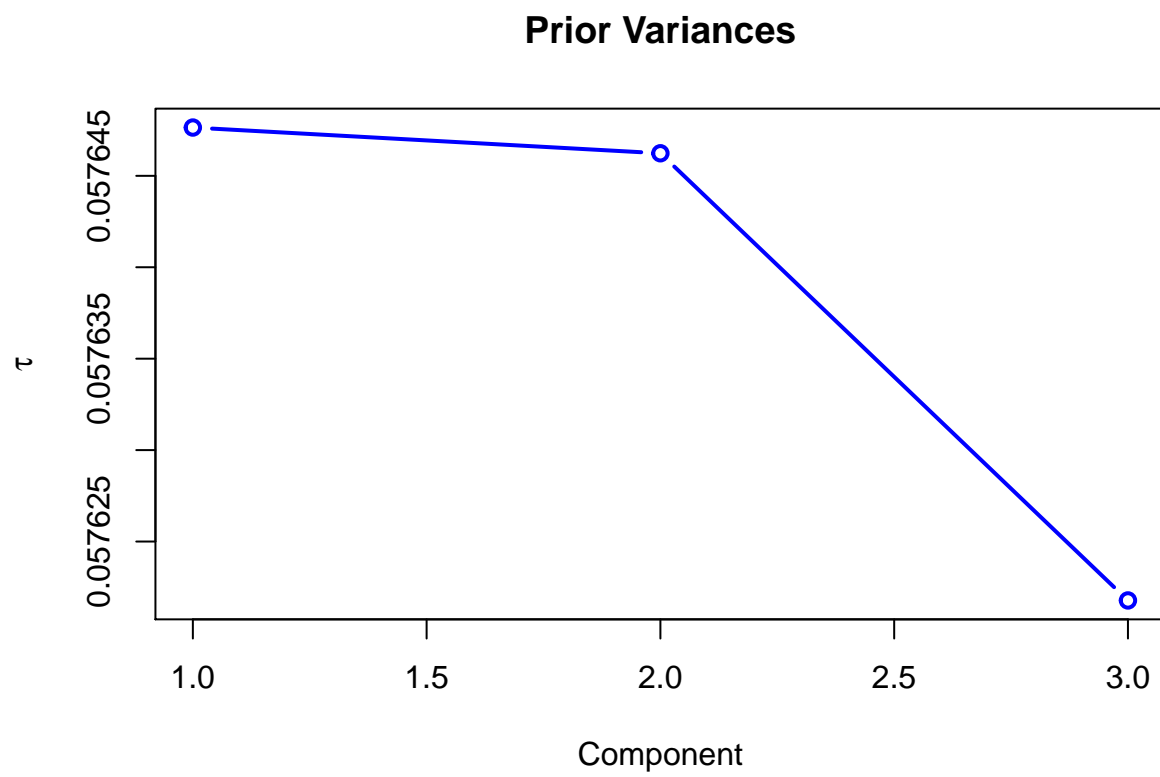


Figure 1: Prior variances for the first 3 components.




Figure 2: Prior variances for 10 components.

```
mod7$invTau
```

```
## [1] 17.34676 17.34719 17.35455
```

Notice the plot of the prior variances (inverse precisions) that appears in this case. This is useful when the number of components supported by the data is uncertain (elbow method):

```
mod8 <- vbPCA(X, D = 10, maxIter = 1e+03, priorvar = 'fixed',
  updatetau = TRUE, global.var = TRUE,
  control = ctrl, verbose = FALSE )
```

Stochastic Search Variable Selection

By requiring `SVS = TRUE`, the model activates stochastic-search-variable-selection, a method described by George and McCulloch (1993) for the Gibbs Sampler. The method has been adapted in *bayespca* for the Variational Bayes algorithm. The assumed ‘spike-and-slab’ prior for the (j, d) -th element of W becomes:

$$w_{j,d} \sim N(0, \pi\tau + (1 - \pi)\tau v_0)$$

where v_0 is a scalar which rescales the spike variance to a value close to 0. For this reason, v_0 should be a number included in $(0, 1)$, as close as possible to 0. π represents the prior probability of inclusion of the j -th variable in the d -th component of the model. `vbPCA` estimates the posterior probabilities of inclusion, conditional on X and the values in W .

While v_0 should be a small value close to 0, too small values of such parameter will shrink the variances τ too much, and no variable will eventually be included in the model. On the other hand, using a too large value for v_0 will not shrink the variances enough, and all posterior inclusion probabilities will be close to 1. v_0 should then be set with a grain of salt. Preliminary results from partial simulation studies have shown that values between 0.0001 and 0.005 lead to acceptable results, but adequate values of v_0 can be dataset-specific. Preliminary simulation studies have also shown that the method works better when Inverse Gamma priors are specified for τ , with **betatau** equal to 1 and larger values of **alphatau**. However, the technique has just been devised and further studies must be carried out to further test the functioning of Stochastic Variable Selection in this context.

In **vbpc**, the parameter v_0 is called **v0** in the control parameters of **vbpc_control**, while the prior inclusion probability is called **priorInclusion**. **priorInclusion** can be fixed, updated via Type-II maximum likelihood, or assigned to a Beta hyperprior:

- among the control parameters of **vbpc_control**, set **beta1pi** smaller than 0 for fixed π ;
- set **beta1pi** = 0 instead for Type-II ML updates;
- last, set **beta1pi** larger than 0 for Beta specifications.

When **beta1pi** is larger than 0, a Beta prior is assumed for π :

$$\pi \sim \text{Beta}(\beta_1, \beta_2).$$

In **vbpc**, β_1 can be controlled with the **beta1pi** argument and β_2 with the **beta2pi** argument in **vbpc_control**.

```
# SVS, fixed priorInclusion and InverseGamma(5, 1) for tau, v0 = .005
ctrl5 <- vbpc_control(center = FALSE, scalecorrection = -1,
                     plot.lowerbound = FALSE,
                     alphatau = 5, betatau = 1,
                     beta1pi = -1, v0 = 5e-03)

# Estimate the model with priorInclusion = 0.5
mod9 <- vbpc(X, D = 3, maxIter = 1e+03, priorvar = 'invgamma',
            SVS = TRUE, priorInclusion = 0.5, control = ctrl5,
            verbose = FALSE )

mod9$muW
```

```
##      Component 1 Component 2 Component 3
## [1,] -0.376748781 -0.04448419 -0.004118464
## [2,] -0.372475633 -0.04361763 -0.005760081
## [3,] -0.375667438 -0.04356739 -0.004991567
## [4,] -0.373841437 -0.04448232 -0.004702291
## [5,] -0.377538204 -0.04319580 -0.005460548
## [6,] -0.375841926 -0.04472555 -0.004526083
## [7,] -0.375415421 -0.04399647 -0.004250611
## [8,]  0.044408685 -0.37703894 -0.019875476
## [9,]  0.043998937 -0.37297946 -0.020143278
## [10,] 0.043266722 -0.37566487 -0.019143767
## [11,] 0.043857102 -0.37129411 -0.019783811
## [12,] 0.044746209 -0.37519130 -0.019756230
## [13,] 0.045169995 -0.37456826 -0.019725427
## [14,] 0.043746355 -0.37732922 -0.018255778
## [15,] 0.002623419  0.02125881 -0.405658519
## [16,] 0.002688735  0.02140113 -0.407985305
## [17,] 0.002848870  0.02135504 -0.409081847
```

```

## [18,] 0.002537808 0.02133019 -0.406355278
## [19,] 0.002543663 0.02127051 -0.408688560
## [20,] 0.002678639 0.02114513 -0.408149455

# SVS, priorInclusion update via Type-II ML and InverseGamma(5, 1) for tau, v0 = .005
ctrl6 <- vbpc_control(center = FALSE, scalecorrection = -1,
                      plot.lowerbound = FALSE,
                      alphatau = 5, betatau = 1,
                      betapi = 0, v0 = 5e-03)

# Estimate the model
mod10 <- vbpc(X, D = 3, maxIter = 1e+03, priorvar = 'invgamma',
              SVS = TRUE, control = ctrl6, verbose = FALSE )

mod10$muW

##           Component 1 Component 2 Component 3
## [1,] -0.376820403 -0.04443274 -0.004105331
## [2,] -0.372342825 -0.04359154 -0.005703954
## [3,] -0.375714669 -0.04343893 -0.004950388
## [4,] -0.373896216 -0.04418282 -0.004672491
## [5,] -0.377494063 -0.04330092 -0.005410152
## [6,] -0.375803977 -0.04448199 -0.004502587
## [7,] -0.375540004 -0.04397016 -0.004233775
## [8,] 0.044272156 -0.37707966 -0.019695158
## [9,] 0.043953350 -0.37295603 -0.019949514
## [10,] 0.043389387 -0.37573980 -0.018994759
## [11,] 0.043792275 -0.37127266 -0.019600074
## [12,] 0.044486863 -0.37516935 -0.019579317
## [13,] 0.044628213 -0.37453668 -0.019545357
## [14,] 0.043972815 -0.37745452 -0.018153343
## [15,] 0.002609130 0.02108194 -0.405657865
## [16,] 0.002665408 0.02120868 -0.407997718
## [17,] 0.002843743 0.02116518 -0.409083064
## [18,] 0.002534612 0.02112394 -0.406366335
## [19,] 0.002532725 0.02108904 -0.408695150
## [20,] 0.002668113 0.02094568 -0.408176915

# SVS, priorInclusion with Beta(1,1) priors and InverseGamma(5, 1) for tau, v0 = .005
ctrl7 <- vbpc_control(center = FALSE, scalecorrection = -1,
                      plot.lowerbound = FALSE, alphatau = 5,
                      betatau = 1, betapi = 1, beta2pi = 1,
                      v0 = 1e-03)

# Estimate the model
mod11 <- vbpc(X, D = 3, maxIter = 1e+03, priorvar = 'invgamma',
              SVS = TRUE, priorInclusion = 0.5, control = ctrl7,
              verbose = FALSE )

mod11$muW

##           Component 1 Component 2 Component 3
## [1,] -0.377578349 -0.047910030 -0.002002242

```

```
## [2,] -0.373170166 -0.042515344 -0.002378816
## [3,] -0.374430126 -0.040415504 -0.002185917
## [4,] -0.375072133 -0.046525265 -0.002135318
## [5,] -0.376008730 -0.039445556 -0.002304975
## [6,] -0.375859884 -0.047111723 -0.002099864
## [7,] -0.375405101 -0.045679887 -0.002032556
## [8,] 0.043405151 -0.378473952 -0.006164428
## [9,] 0.044375902 -0.373876173 -0.006209141
## [10,] 0.042245921 -0.377090189 -0.006001771
## [11,] 0.044751842 -0.368672388 -0.006116497
## [12,] 0.044958898 -0.376869033 -0.006128563
## [13,] 0.045638798 -0.375085381 -0.006120630
## [14,] 0.044350991 -0.377236487 -0.005819845
## [15,] 0.001437227 0.006323692 -0.405995608
## [16,] 0.001446743 0.006322625 -0.408537375
## [17,] 0.001465363 0.006378243 -0.409385330
## [18,] 0.001440418 0.006341427 -0.406803848
## [19,] 0.001424396 0.006367079 -0.409223708
## [20,] 0.001427028 0.006333777 -0.409086743
```

The estimated posterior inclusion probabilities for the three models:

```
mod9$inclusionProbabilities
```

```
##           [,1]      [,2]      [,3]
## [1,] 1.00000000 0.2119197 0.09722059
## [2,] 1.00000000 0.2067622 0.09847080
## [3,] 1.00000000 0.2083132 0.09893331
## [4,] 1.00000000 0.2138161 0.09839703
## [5,] 1.00000000 0.2043346 0.09818925
## [6,] 1.00000000 0.2145334 0.09844897
## [7,] 1.00000000 0.2089924 0.09789326
## [8,] 0.21419852 1.0000000 0.11720313
## [9,] 0.21103005 1.0000000 0.11749524
## [10,] 0.20602971 1.0000000 0.11538121
## [11,] 0.21062206 1.0000000 0.11704114
## [12,] 0.21585734 1.0000000 0.11699144
## [13,] 0.21973753 1.0000000 0.11755002
## [14,] 0.20652527 1.0000000 0.11220954
## [15,] 0.09656056 0.1178761 1.00000000
## [16,] 0.09698534 0.1193360 1.00000000
## [17,] 0.09664596 0.1184619 1.00000000
## [18,] 0.09701643 0.1191947 1.00000000
## [19,] 0.09630871 0.1177812 1.00000000
## [20,] 0.09738379 0.1190830 1.00000000
```

```
mod10$inclusionProbabilities
```

```
##           [,1]      [,2]      [,3]
## [1,] 1.00000000 0.14595571 0.06591143
## [2,] 1.00000000 0.14236100 0.06674450
## [3,] 1.00000000 0.14281308 0.06703891
## [4,] 1.00000000 0.14589407 0.06668594
## [5,] 1.00000000 0.14122352 0.06655497
## [6,] 1.00000000 0.14674973 0.06671996
## [7,] 1.00000000 0.14397192 0.06635530
```

```
## [8,] 0.14688020 1.00000000 0.07946508
## [9,] 0.14513342 1.00000000 0.07965978
## [10,] 0.14241065 1.00000000 0.07827022
## [11,] 0.14472254 1.00000000 0.07935064
## [12,] 0.14747139 1.00000000 0.07932312
## [13,] 0.14867799 1.00000000 0.07967706
## [14,] 0.14348255 1.00000000 0.07619741
## [15,] 0.06545289 0.08000438 1.00000000
## [16,] 0.06572667 0.08094135 1.00000000
## [17,] 0.06550593 0.08036715 1.00000000
## [18,] 0.06574893 0.08082310 1.00000000
## [19,] 0.06528540 0.07993301 1.00000000
## [20,] 0.06598849 0.08074012 1.00000000
```

```
mod11$inclusionProbabilities
```

```
##           [,1]      [,2]      [,3]
## [1,] 1.00000000 1.00000000 0.06863487
## [2,] 1.00000000 1.00000000 0.06904590
## [3,] 1.00000000 1.00000000 0.06902865
## [4,] 1.00000000 1.00000000 0.06891943
## [5,] 1.00000000 1.00000000 0.06891337
## [6,] 1.00000000 1.00000000 0.06891772
## [7,] 1.00000000 1.00000000 0.06880151
## [8,] 1.00000000 1.00000000 0.07472949
## [9,] 1.00000000 1.00000000 0.07482446
## [10,] 1.00000000 1.00000000 0.07440630
## [11,] 1.00000000 1.00000000 0.07469438
## [12,] 1.00000000 1.00000000 0.07466636
## [13,] 1.00000000 1.00000000 0.07472689
## [14,] 1.00000000 1.00000000 0.07375244
## [15,] 0.06813425 0.07458764 1.00000000
## [16,] 0.06826714 0.07472781 1.00000000
## [17,] 0.06805943 0.07461328 1.00000000
## [18,] 0.06826063 0.07476505 1.00000000
## [19,] 0.06803569 0.07458175 1.00000000
## [20,] 0.06822868 0.07472175 1.00000000
```

It is also possible to compare the (known) variable inclusion matrix vs. the estimated ones graphically. Let's plot a heatmap of such probabilities for model mod9:

```
trueInclusions <- matrix(0, J, 3)
trueInclusions[1:7, 1] <- 1
trueInclusions[8:14, 2] <- 1
trueInclusions[15:20, 3] <- 1

par(mfrow=c(1,2))
image(1:ncol(trueInclusions), 1:nrow(trueInclusions),
      t(trueInclusions[J:1, ]), ylab = "", axes = FALSE,
      main = "True Inclusions", xlab = "",
      col = RColorBrewer::brewer.pal(9, "Blues"))
axis(side = 1, at = 1:3, labels = paste("Component ", 1:3 ))
axis(side = 2, at = 1:20, labels = paste("Var ", J:1 ))

fields::image.plot(1:ncol(trueInclusions), 1:nrow(trueInclusions),
                   t(mod9$inclusionProbabilities[J:1, ]), ylab = "", axes = FALSE,
```

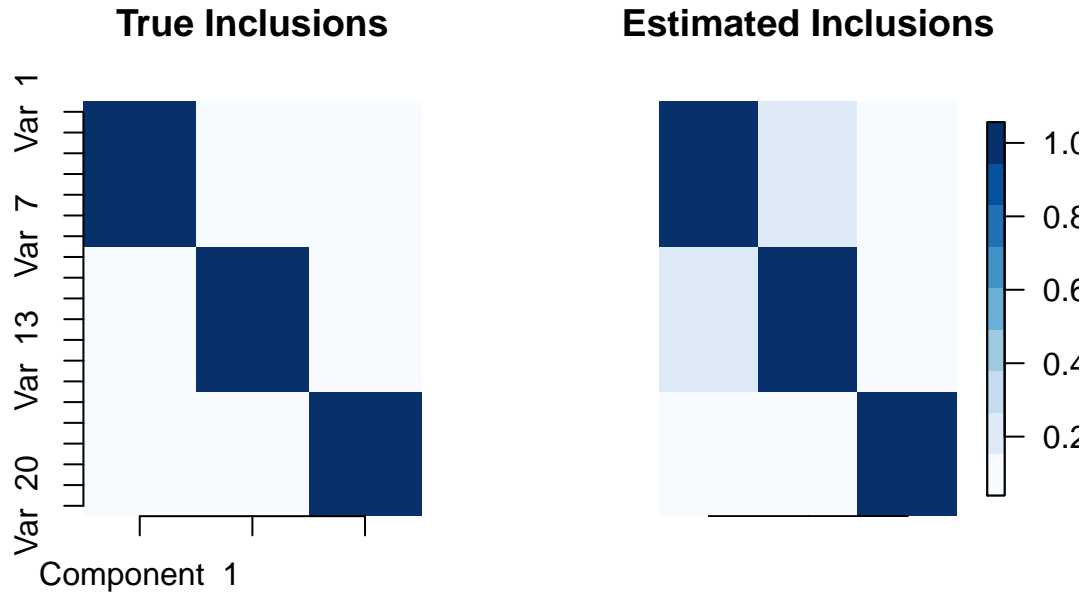



Figure 3: True and Estimated inclusion probabilities.

```
main = "Estimated Inclusions", xlab = "",
col = RColorBrewer::brewer.pal(9, "Blues"))
axis(side = 1, at = 1:3, labels = paste("Component ", 1:3 ))
```

For mod10 and mod11 we can observe the estimated prior inclusion probabilities:

```
mod10$priorInclusion
```

```
##           [,1]
## [1,] 0.3987037
## [2,] 0.3987037
## [3,] 0.3987037
```

```
mod11$priorInclusion
```

```
##           [,1]
## [1,] 0.594533
## [2,] 0.594533
## [3,] 0.594533
```

Similar to the hyperparameters of the Inverse Gamma priors on τ , `priorInclusion`, `beta1pi` and `beta2pican` also be specified as D -dimensional arrays. This will allow estimating the inclusion probabilities with different degrees of ‘sparsity’ for each component. For Type-II maximum likelihood updates of `priorInclusion`, all elements of `beta1pi` must be equal to 0. For Beta priors, all elements of `beta1pi` must be larger than 0. Let us look at two examples:

```
# Type-II maximum likelihood (component-specific) updates
ctrl18 <- vbpcr_control(center = FALSE, scalecorrection = -1,
```

```

        plot.lowerbound = FALSE,
        alphatau = 5, betatau = 1,
        betapi = rep(0, 3), v0 = 5e-03)

# Estimate the model
mod12 <- vbPCA(X, D = 3, maxIter = 1e+03, priorvar = 'invgamma',
              SVS = TRUE, priorInclusion = rep(0.5, 3),
              control = ctrl8, verbose = FALSE )

mod12$muW

```

```

##           Component 1 Component 2 Component 3
## [1,] -0.376815314 -0.04442994 -0.004105677
## [2,] -0.372377110 -0.04356041 -0.005685535
## [3,] -0.375706910 -0.04340680 -0.004938422
## [4,] -0.373904118 -0.04418880 -0.004665602
## [5,] -0.377491770 -0.04324636 -0.005394055
## [6,] -0.375805352 -0.04449342 -0.004498355
## [7,] -0.375526972 -0.04394678 -0.004232412
## [8,]  0.044263800 -0.37707627 -0.019646180
## [9,]  0.043932375 -0.37297930 -0.019894678
## [10,] 0.043343358 -0.37571010 -0.018958464
## [11,] 0.043768765 -0.37129276 -0.019549524
## [12,] 0.044484539 -0.37519022 -0.019531727
## [13,] 0.044640336 -0.37456687 -0.019496815
## [14,] 0.043933453 -0.37742346 -0.018136623
## [15,] 0.002608240  0.02103643 -0.405655232
## [16,] 0.002665114  0.02116643 -0.408001312
## [17,] 0.002841361  0.02111936 -0.409083269
## [18,] 0.002533538  0.02108528 -0.406367096
## [19,] 0.002530947  0.02104353 -0.408696279
## [20,] 0.002665016  0.02090398 -0.408188568

```

```
mod12$priorInclusion
```

```

##           [,1]
## [1,] 0.4294228
## [2,] 0.4359323
## [3,] 0.3399249

```

```
mod12$inclusionProbabilities
```

```

##           [,1]      [,2]      [,3]
## [1,] 1.00000000 0.16799387 0.05168586
## [2,] 1.00000000 0.16374569 0.05233463
## [3,] 1.00000000 0.16430571 0.05255872
## [4,] 1.00000000 0.16801270 0.05228552
## [5,] 1.00000000 0.16234229 0.05218633
## [6,] 1.00000000 0.16899906 0.05231170
## [7,] 1.00000000 0.16562408 0.05202976
## [8,] 0.16501213 1.00000000 0.06235012
## [9,] 0.16298497 1.00000000 0.06250196
## [10,] 0.15981506 1.00000000 0.06142445
## [11,] 0.16252283 1.00000000 0.06225847

```

```

## [12,] 0.16569009 1.00000000 0.06223902
## [13,] 0.16713070 1.00000000 0.06250883
## [14,] 0.16099504 1.00000000 0.05982298
## [15,] 0.07386198 0.09239017 1.00000000
## [16,] 0.07417472 0.09348813 1.00000000
## [17,] 0.07392333 0.09281332 1.00000000
## [18,] 0.07420049 0.09335883 1.00000000
## [19,] 0.07367201 0.09230762 1.00000000
## [20,] 0.07447276 0.09326539 1.00000000

# Beta priors with different degrees of sparsity for each component
ctrl9 <- vbPCA_control(center = FALSE, scalecorrection = -1,
                        plot.lowerbound = FALSE,
                        alphatau = 5, betatau = 1,
                        beta1pi = c(0.01, 1, 10), beta2pi = 1,
                        v0 = 5e-03)

# Estimate the model
mod13 <- vbPCA(X, D = 3, maxIter = 1e+03, priorvar = 'invgamma', SVS = TRUE,
               priorInclusion = rep(0.5, 3), control = ctrl9, verbose = FALSE )

mod13$muW

##           Component 1 Component 2 Component 3
## [1,] -0.376821096 -0.04445239 -0.004064560
## [2,] -0.372338748 -0.04357489 -0.005758008
## [3,] -0.375715573 -0.04342516 -0.004973482
## [4,] -0.373895493 -0.04421742 -0.004668698
## [5,] -0.377494191 -0.04325598 -0.005453514
## [6,] -0.375803733 -0.04452144 -0.004484911
## [7,] -0.375541802 -0.04396579 -0.004201890
## [8,]  0.044273060 -0.37707142 -0.019747340
## [9,]  0.043955743 -0.37297355 -0.020032637
## [10,] 0.043394788 -0.37570751 -0.018982867
## [11,] 0.043794756 -0.37128099 -0.019663100
## [12,] 0.044486601 -0.37518654 -0.019625116
## [13,] 0.044624975 -0.37456067 -0.019597277
## [14,] 0.043977647 -0.37741749 -0.018039117
## [15,] 0.002609693  0.02110669 -0.405684212
## [16,] 0.002665877  0.02123764 -0.407991248
## [17,] 0.002844509  0.02119104 -0.409091468
## [18,] 0.002535224  0.02115675 -0.406376351
## [19,] 0.002533398  0.02111446 -0.408701144
## [20,] 0.002668911  0.02097557 -0.408126495

mod13$priorInclusion

##           [,1]
## [1,] 0.3999238
## [2,] 0.4431880
## [3,] 0.5825101

mod13$inclusionProbabilities

##           [,1]      [,2]      [,3]

```

```
## [1,] 1.00000000 0.17100954 0.1336448
## [2,] 1.00000000 0.16664977 0.1353853
## [3,] 1.00000000 0.16724781 0.1360785
## [4,] 1.00000000 0.17107270 0.1353160
## [5,] 1.00000000 0.16519843 0.1349989
## [6,] 1.00000000 0.17206838 0.1353922
## [7,] 1.00000000 0.16858209 0.1345997
## [8,] 0.14480895 1.00000000 0.1602183
## [9,] 0.14309364 1.00000000 0.1606159
## [10,] 0.14042124 1.00000000 0.1576725
## [11,] 0.14268762 1.00000000 0.1600187
## [12,] 0.14538795 1.00000000 0.1599328
## [13,] 0.14656203 1.00000000 0.1607567
## [14,] 0.14148141 1.00000000 0.1531830
## [15,] 0.06449702 0.09415751 1.0000000
## [16,] 0.06476642 0.09528034 1.0000000
## [17,] 0.06454915 0.09459244 1.0000000
## [18,] 0.06478830 0.09514877 1.0000000
## [19,] 0.06433207 0.09407460 1.0000000
## [20,] 0.06502416 0.09505391 1.0000000
```

High posterior density intervals

It is also possible to require the computation of high probability density intervals for the elements of W , which can then be plotted with the `plothpdi` function, which internally calls `ggplot2` functionalities.

```
# Set hyperparameter values and require 50% probability density intervals
ctrl10 <- vbpca_control(center = FALSE, scalecorrection = -1,
                        plot.lowerbound = FALSE,
                        alphatau = 2, betatau = .5,
                        hpdi = TRUE, probHPDI = 0.5)

# Estimate the model
mod14 <- vbpca(X, D = 3, maxIter = 1e+03, priorvar = 'invgamma',
               control = ctrl10, verbose = TRUE )
```

```
## Local prior variances : Inverse-Gamma, fixed hyperparameters.
## Start # 1 has converged in 2 iterations; lower bound = -2802.92

# Plot HPD intervals for variables 1:10 component 1
plothpdi(mod14, d = 1, X, vars = 1:10)
```

Retrieve Principal Components

To compute the estimated components, simply call:

```
PCs <- X %*% mod1$muW
head(PCs, 15)
```

```
##      Component 1 Component 2 Component 3
## [1,]   -59.19132   -78.592707   31.3401056
## [2,]    28.97173  -118.789002  -29.0200803
## [3,]   -11.00518    14.227039   -4.8429367
## [4,]    92.16140   -33.606390  -28.1184509
```




Figure 4: High posterior density intervals.

```
## [5,] -41.61482 -212.440559 13.4800647
## [6,] 113.51610 -20.107248 5.6778548
## [7,] 98.45308 -73.892683 17.2711826
## [8,] 42.05467 -142.922658 -68.0937551
## [9,] -57.38540 -66.586047 17.5396918
## [10,] 42.94090 51.286634 -0.2553017
## [11,] 36.39523 -11.871548 13.9383095
## [12,] 109.60474 -6.656482 25.3900580
## [13,] -196.01791 110.020825 -9.5996919
## [14,] -267.42318 71.336729 14.1676697
## [15,] 38.49334 22.034659 -32.6994089
```

References

1. C. M. Bishop. ‘Variational PCA’. In Proc. Ninth Int. Conf. on Artificial Neural Networks. ICANN, 1999.
2. E. I. George, R. E. McCulloch (1993). ‘Variable Selection via Gibbs Sampling’. Journal of the American Statistical Association (88), 881-889.